

Raft lab

• Software

You'll implement this lab (and all the labs) in [Go](#). The Go web site contains lots of tutorial information. We will grade your labs using Go version 1.13; you should use 1.13 too. You can check your Go version by running `go version`.

We recommend that you work on the labs on your own machine, so you can use the tools, text editors, etc. that you are already familiar with.

We suggest download directly from [Go](#) and follow the steps. Make sure that you create your workspace directory (see more information in [golang.org](#)).

◦ Windows

The labs probably won't work directly on Windows. If you're feeling adventurous, you can try to get them running inside [Windows Subsystem for Linux](#).

NOTICE: Do not put the lab in your \$GOPATH dir.

Test everything is ready:

```
1 go version // make sure it's 1.13.x
2 cd raft_lab/src/raft
3 go test
4 Test (2A): initial election ...
5 --- FAIL: TestInitialElection2A (5.04s)
6         config.go:326: expected one leader, got none
7 Test (2A): election after network failure ...
8 --- FAIL: TestReElection2A (5.03s)
9         config.go:326: expected one leader, got none
10 ...
```

• Introduction

We supply you with skeleton code `src/raft/raft.go`. We also supply a set of tests, which you should use to drive your implementation efforts, and which we'll use to grade your submitted lab. The tests are in `src/raft/test_test.go`.

• The code

Implement Raft by adding code to `raft/raft.go`

Your implementation must support the following interface, which the tester and (eventually) your key/value server will use. You'll find more details in comments in `raft.go`.

```

1 // create a new Raft server instance:
2 rf := Make(peers, me, persister, applyCh)
3
4 // start agreement on a new log entry:
5 rf.Start(command interface{}) (index, term, isleader)
6
7 // ask a Raft for its current term, and whether it thinks it is leader
8 rf.GetState() (term, isLeader)
9
10 // each time a new entry is committed to the log, each Raft peer
11 // should send an ApplyMsg to the service (or tester).
12 type ApplyMsg

```

A service calls `Make(peers, me, ...)` to create a Raft peer. The `peers` argument is an array of network identifiers of the Raft peers (including this one), for use with RPC. The `me` argument is the index of this peer in the `peers` array. `Start(command)` asks Raft to start the processing to append the command to the replicated log. `Start()` should return immediately, without waiting for the log appends to complete. The service expects your implementation to send an `ApplyMsg` for each newly committed log entry to the `applyCh` channel argument to `Make()`.

`raft.go` contains example code that sends an RPC (`sendRequestVote()`) and that handles an incoming RPC (`RequestVote()`). Your Raft peers should exchange RPCs using the `labrpc` Go package (source in `src/labrpc`). The tester can tell `labrpc` to delay RPCs, re-order them, and discard them to simulate various network failures. While you can temporarily modify `labrpc`, make sure your Raft works with the original `labrpc`, since that's what we'll use to test and grade your lab. Your Raft instances must interact only with RPC; for example, they are not allowed to communicate using shared Go variables or files. We will check this.

Subsequent labs build on this lab, so it is important to give yourself enough time to write solid code.

• Part A

TASK:

Implement Raft leader election and heartbeats (`AppendEntries` RPCs with no log entries). The goal for Part 2A is for a single leader to be elected, for the leader to remain the leader if there are no failures, and for a new leader to take over if the old leader fails or if packets to/from the old leader are lost. Run `go test -run 2A` to test your 2A code.

Some hints:

1. You can't easily run your Raft implementation directly; instead you should run it by way of the tester, i.e. `go test -run 2A`.
2. Follow the paper's Figure 2. At this point you care about sending and receiving `RequestVote` RPCs, the Rules for Servers that relate to elections, and the State related to leader election.
3. Add the Figure 2 state for leader election to the `Raft` struct in `raft.go`. You'll also need to define a struct to hold information about each log entry.

4. Fill in the `RequestVoteArgs` and `RequestVoteReply` structs. Modify `Make()` to create a background goroutine that will kick off leader election periodically by sending out `RequestVote` RPCs when it hasn't heard from another peer for a while. This way a peer will learn who is the leader, if there is already a leader, or become the leader itself. Implement the `RequestVote()` RPC handler so that servers will vote for one another.
5. To implement heartbeats, define an `AppendEntries` RPC struct (though you may not need all the arguments yet), and have the leader send them out periodically. Write an `AppendEntries` RPC handler method that resets the election timeout so that other servers don't step forward as leaders when one has already been elected.
6. Make sure the election timeouts in different peers don't always fire at the same time, or else all peers will vote only for themselves and no one will become the leader.
7. The tester requires that the leader send heartbeat RPCs no more than ten times per second.
8. The tester requires your Raft to elect a new leader within five seconds of the failure of the old leader (if a majority of peers can still communicate). Remember, however, that leader election may require multiple rounds in case of a split vote (which can happen if packets are lost or if candidates unluckily choose the same random backoff times). You must pick election timeouts (and thus heartbeat intervals) that are short enough that it's very likely that an election will complete in less than five seconds even if it requires multiple rounds.
9. The paper's Section 5.2 mentions election timeouts in the range of 150 to 300 milliseconds. Such a range only makes sense if the leader sends heartbeats considerably more often than once per 150 milliseconds. Because the tester limits you to 10 heartbeats per second, you will have to use an election timeout larger than the paper's 150 to 300 milliseconds, but not too large, because then you may fail to elect a leader within five seconds.
10. You may find Go's [rand](#) useful.
11. You'll need to write code that takes actions periodically or after delays in time. The easiest way to do this is to create a goroutine with a loop that calls [time.Sleep\(\)](#). Don't use Go's `time.Timer` or `time.Ticker`, which are difficult to use correctly.
12. Read this advice about locking([raft-locking.txt](#)) and structure([raft-structure.txt](#)).
13. If your code has trouble passing the tests, read the paper's Figure 2 again; the full logic for leader election is spread over multiple parts of the figure.
14. Don't forget to implement `GetState()`.
15. The tester calls your Raft's `rf.Kill()` when it is permanently shutting down an instance. You can check whether `Kill()` has been called using `rf.killed()`. You may want to do this in all loops, to avoid having dead Raft instances print confusing messages.
16. A good way to debug your code is to insert print statements when a peer sends or receives a message, and collect the output in a file with `go test -run 2A > out`. Then, by studying the trace of messages in the `out` file, you can identify where your implementation deviates from the desired protocol. You might find `DPrintf` in `util.go` useful to turn printing on and off as you debug different problems.
17. Go RPC sends only struct fields whose names start with capital letters. Sub-structures must also have capitalized field names (e.g. fields of log records in an array). The `labgob` package will warn you about this; don't ignore the warnings.

18. Check your code with `go test -race`, and fix any races it reports.
19. Run a single test(e.g. `TestElection2A`) with `go test -run Election2A`

Be sure you pass the 2A tests before submitting Part 2A, so that you see something like this:

```
1 $ go test -run 2A
2 Test (2A): initial election ...
3 ... Passed -- 4.0 3 32 9170 0
4 Test (2A): election after network failure ...
5 ... Passed -- 6.1 3 70 13895 0
6 PASS
7 ok      raft    10.187s
8 $
```

Each "Passed" line contains five numbers; these are the time that the test took in seconds, the number of Raft peers (usually 3 or 5), the number of RPCs sent during the test, the total number of bytes in the RPC messages, and the number of log entries that Raft reports were committed. Your numbers will differ from those shown here. You can ignore the numbers if you like, but they may help you sanity-check the number of RPCs that your implementation sends. For all of labs 2, 3, and 4, the grading script will fail your solution if it takes more than 600 seconds for all of the tests (`go test`), or if any individual test takes more than 120 seconds.

• Part B

TASK:

Implement the leader and follower code to append new log entries, so that the `go test -run 2B` tests pass.

Some hints:

1. Your first goal should be to pass `TestBasicAgree2B()`. Start by implementing `start()`, then write the code to send and receive new log entries via `AppendEntries` RPCs, following Figure 2.
2. You will need to implement the election restriction (section 5.4.1 in the paper).
3. One way to fail to reach agreement in the early Lab 2B tests is to hold repeated elections even though the leader is alive. Look for bugs in election timer management, or not sending out heartbeats immediately after winning an election.
4. Your code may have loops that repeatedly check for certain events. Don't have these loops execute continuously without pausing, since that will slow your implementation enough that it fails tests. Use Go's condition variables, or insert a `time.Sleep(10 * time.Millisecond)` in each loop iteration.
5. Do yourself a favor for future labs and write (or re-write) code that's clean and clear.

The tests for upcoming labs may fail your code if it runs too slowly. You can check how much real time and CPU time your solution uses with the `time` command. Here's typical output:

```

1  $ time go test -run 2B
2  Test (2B): basic agreement ...
3    ... Passed --    1.6  3   18   5158    3
4  Test (2B): RPC byte count ...
5    ... Passed --    3.3  3   50  115122   11
6  Test (2B): agreement despite follower disconnection ...
7    ... Passed --    6.3  3   64   17489    7
8  Test (2B): no agreement if too many followers disconnect ...
9    ... Passed --    4.9  5  116   27838    3
10 Test (2B): concurrent Start()s ...
11    ... Passed --    2.1  3   16   4648    6
12 Test (2B): rejoin of partitioned leader ...
13    ... Passed --    8.1  3  111   26996    4
14 Test (2B): leader backs up quickly over incorrect follower logs ...
15    ... Passed --   28.6  5 1342  953354  102
16 Test (2B): RPC counts aren't too high ...
17    ... Passed --    3.4  3   30   9050   12
18 PASS
19 ok      raft      58.142s
20
21 real    0m58.475s
22 user    0m2.477s
23 sys     0m1.406s
24 $

```

The "ok raft 58.142s" means that Go measured the time taken for the 2B tests to be 58.142 seconds of real (wall-clock) time. The "user 0m2.477s" means that the code consumed 2.477 seconds of CPU time, or time spent actually executing instructions (rather than waiting or sleeping). If your solution uses much more than a minute of real time for the 2B tests, or much more than 5 seconds of CPU time, you may run into trouble later on. Look for time spent sleeping or waiting for RPC timeouts, loops that run without sleeping or waiting for conditions or channel messages, or large numbers of RPCs sent.

• Part C

If a Raft-based server reboots it should resume service where it left off. This requires that Raft keep persistent state that survives a reboot. The paper's Figure 2 mentions which state should be persistent.

A real implementation would write Raft's persistent state to disk each time it changed, and would read the state from disk when restarting after a reboot. Your implementation won't use the disk; instead, it will save and restore persistent state from a `Persister` object (see `persister.go`). Whoever calls `Raft.Make()` supplies a `Persister` that initially holds Raft's most recently persisted state (if any). Raft should initialize its state from that `Persister`, and should use it to save its persistent state each time the state changes. Use the `Persister`'s `ReadRaftState()` and `SaveRaftState()` methods.

Task:

1. Complete the functions `persist()` and `readPersist()` in `raft.go` by adding code

to save and restore persistent state. You will need to encode (or "serialize") the state as an array of bytes in order to pass it to the `Persister`. Use the `labgob` encoder; see the comments in `persist()` and `readPersist()`. `labgob` is like Go's `gob` encoder but prints error messages if you try to encode structures with lower-case field names.

2. Insert calls to `persist()` at the points where your implementation changes persistent state. Once you've done this, you should pass the remaining tests.

Note:

In order to avoid running out of memory, Raft must periodically discard old log entries, but you **do not** have to worry about this.

Hints:

1. Many of the 2C tests involve servers failing and the network losing RPC requests or replies.
2. You will probably need the optimization that backs up `nextIndex` by more than one entry at a time. Look at the Raft paper starting at the bottom of page 7 and top of page 8 (marked by a gray line). The paper is vague about the details; you will need to fill in the gaps.
3. A reasonable amount of time to consume for the full set of Lab 2 tests (2A+2B+2C) is 4 minutes of real time and one minute of CPU time.

Your code should pass all the 2C tests (as shown below), as well as the 2A and 2B tests.

```
1  $ go test -run 2C
2  Test (2C): basic persistence ...
3    ... Passed --   7.2  3  206  42208    6
4  Test (2C): more persistence ...
5    ... Passed --  23.2  5 1194 198270   16
6  Test (2C): partitioned leader and one follower crash, leader restarts
7    ...
8    ... Passed --   3.2  3   46  10638    4
9  Test (2C): Figure 8 ...
10   ... Passed --  35.1  5 9395 1939183   25
11 Test (2C): unreliable agreement ...
12   ... Passed --   4.2  5  244  85259   246
13 Test (2C): Figure 8 (unreliable) ...
14   ... Passed --  36.3  5 1948 4175577   216
15 Test (2C): churn ...
16   ... Passed --  16.6  5 4402 2220926 1766
17 Test (2C): unreliable churn ...
18   ... Passed --  16.5  5  781  539084   221
19 PASS
20 ok      raft      142.357s
21 $
```

NOTICE:

Run **ALL** tests together and run multi times. Some mistake will not appear during one running.